

01 - Supervised Learning - Classification - AdaBoost(Solution)_st122097_thantham

September 20, 2021

1 Programming for Data Science and Artificial Intelligence

1.1 Classification - AdaBoost

1.2 Name: Thantham Khamyai

1.3 Student ID: 122097

1.3.1 ===Task===

Your work: Let's modify the above scratch code: - Notice that if $\text{err} = 0$, then α will be undefined, thus attempt to fix this by adding some very small value to the lower term - Notice that sklearn version of AdaBoost has a parameter `learning_rate`. This is in fact the $\frac{1}{2}$ in front of the α calculation. Attempt to change this $\frac{1}{2}$ into a parameter called `eta`, and try different values of it and see whether accuracy is improved. Note that sklearn default this value to 1. - Observe that we are actually using sklearn `DecisionTreeClassifier`. If we take a look at it closely, it is actually using weighted gini index, instead of weighted errors that we learn above. Attempt to write your own class of class `Stump` that actually uses weighted errors, instead of weighted gini index - Put everything into a class

```
[1]: from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
import numpy as np

X, y = make_classification(n_samples=500, random_state=48)
y = np.where(y==0,-1,1) #change our y to be -1 if it is 0, otherwise 1

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42)
```


Show the scatter of dataset before for imagining the dataset looks like. However, the dataset was 20 dimension features, we therefore just show it simply first and last feature combination

```
[2]: import matplotlib.pyplot as plt

plt.scatter(X[:,0], X[:, -1], c=y)
plt.title('Feature space of first and last feature')
plt.xlabel('First feature')
```

```
plt.ylabel('Last feature')
plt.grid()
plt.show()

print('Shape of dataset: ',X.shape)
```



Shape of dataset: (500, 20)

1.3.2 TASK 3: Stump class

Actually, we have to put stump class may be into inner class, but we personally have separated class called **StumpClassifier** to further use separately for different approach.

This is also make the **Adaboost** class being more feasible to be used for different kinds of classifier

```
[3]: class StumpClassifier:

    def __init__(self):
        self.polarity = 1
        self.feature_index = None
        self.threshold = None

    def fit(self, X, y):
```

```

self.min_err = np.inf
m,n = X.shape
W = np.full(m, 1/m) # init equal weight

for feature in range(n): # looping for all features

    feature_vals = np.sort(np.unique(X[:, feature])) # get sorted
    ↪ features
    thresholds = (feature_vals[:-1] + feature_vals[1:])/2 # get all
    ↪ of thresholds in features

    for threshold in thresholds: # looping for each threshold

        for polarity in [1, -1]: # check for polarity twice

            yhat = np.ones(len(y)) # init all answers as 1

            # change the answer into -1 if the feature value that
            ↪ less than threshold in current polarity
            yhat[polarity * X[:, feature] < polarity * threshold] =
            ↪ -1

            # evaluate error
            err = W[(yhat != y)].sum()

            # if error is less than before -> define current
            ↪ feature_idx and threshold of splitting
            if err < self.min_err:
                self.polarity = polarity
                self.threshold = threshold
                self.feature_index = feature
                self.min_err = err

            # if further feature splitting gives more less of error ->
            ↪ define that feature for feature split

def predict(self, X):
    m, n = X.shape

    pred = np.ones(m) # init all answer as 1

    # change the answer into -1 if the feature value that less than
    ↪ threshold in current polarity

```

```

        pred[self.polarity * X[:,self.feature_index] < self.polarity * self.
→threshold] = -1

        return np.sign(pred) # return sign of the answer

```

1.3.3 TASK 4: Put every thing into a class

I made this class similar to the way of sklearn do. we can put any classifier into the this class.

```

[4]: from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC

class AdaBoost:

    def __init__(self, estimator = DecisionTreeClassifier, n_estimators=20,
→model_params={}, eta=0.5):

        self.n_estimators = n_estimators
        self.model_params = model_params
        self.eps = 1e-10 # < ===== TASK 1: Epsilon
        self.eta = eta # < ===== TASK 2: Adding learning rate
→(ETA)

        # init estimator classifiers set following the predefined params
        self.models = [estimator(**self.model_params) for _ in range(self.
→n_estimators)]

    def fit(self, X_train, y_train):

        m, n = X_train.shape

        self.W = np.full(m, 1/m) # init weak learner weights of samples

        self.alpha = np.zeros(self.n_estimators) # init alpha of each classifier

        for i, model in enumerate(self.models): # looping each model in modelset

            model.fit(X_train, y_train) # simply fit
            yhat = model.predict(X_train) #simply predict

        err = self.W[(y_train != yhat)].sum() # evaluate error of current model

        # calculate alpha (voting power) of current model

```

```

        self.alpha = self.eta * np.log((1-err)/(err + self.eps)) # < == Eps
    ↪ implementation

    # calculate Weaker learner weights of each samples
    self.W = (self.W * np.exp(-self.alpha*yhat*y_train) ) / np.sum(self.W)

def predict(self, X_test):

    H_X = 0

    for i, model in enumerate(self.models): # looping for each model in
    ↪ modelset

        yhat = model.predict(X_test) # simply predict
        H_X += self.alpha * yhat # weight the answers

    return np.sign(H_X) # just give the sign of answers

```

To show up the model class, I put **StumpClassifier** for estimator and keep it default number of estimators

```

[5]: model = AdaBoost(estimator=StumpClassifier, n_estimators=20, model_params={},
    ↪ eta=0.5)
    model.fit(X_train, y_train)

```

As we known, lets predict the result of training

```

[6]: y_pred = model.predict(X_test)

```

The result is also being the same as set of $\{-1, +1\}$ following therotical

```

[7]: y_pred[:20]

```

```

[7]: array([ 1.,  1.,  1., -1.,  1., -1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,
          1., -1.,  1.,  1., -1., -1.,  1.])

```

Finally, examine the classification report

```

[8]: from sklearn.metrics import classification_report

    print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
-1	0.80	0.80	0.80	70
1	0.82	0.82	0.82	80
accuracy			0.81	150

macro avg	0.81	0.81	0.81	150
weighted avg	0.81	0.81	0.81	150

1.3.4 Adaboost Alternative Usage

Lets try implementing different way of **Adaboost**: KNN, SVM, DecisionTree

```
[9]: ada_knn = AdaBoost(estimator=KNeighborsClassifier, n_estimators=10,
    ↪model_params={'n_neighbors': 5}, eta=0.5)
ada_knn.fit(X_train, y_train)
print(classification_report(y_test, ada_knn.predict(X_test)))
```

	precision	recall	f1-score	support
-1	0.77	0.80	0.78	70
1	0.82	0.79	0.80	80
accuracy			0.79	150
macro avg	0.79	0.79	0.79	150
weighted avg	0.79	0.79	0.79	150

```
[10]: ada_svm = AdaBoost(estimator=SVC, n_estimators=5, model_params={}, eta=0.5)
ada_svm.fit(X_train, y_train)
print(classification_report(y_test, ada_svm.predict(X_test)))
```

	precision	recall	f1-score	support
-1	0.76	0.81	0.79	70
1	0.83	0.78	0.80	80
accuracy			0.79	150
macro avg	0.79	0.79	0.79	150
weighted avg	0.80	0.79	0.79	150

```
[11]: ada_dt = AdaBoost(estimator=DecisionTreeClassifier, n_estimators=5,
    ↪model_params={'max_depth': 3}, eta=0.5)
ada_dt.fit(X_train, y_train)
print(classification_report(y_test, ada_dt.predict(X_test)))
```

	precision	recall	f1-score	support
-1	0.73	0.87	0.80	70
1	0.87	0.72	0.79	80
accuracy			0.79	150
macro avg	0.80	0.80	0.79	150

weighted avg	0.80	0.79	0.79	150
--------------	------	------	------	-----

StumpClassifier Usage For show up the usage of Stump classifier

```
[12]: stump = StumpClassifier()
      stump.fit(X_train, y_train)
      print(classification_report(y_test, stump.predict(X_test)))

      print('Feature to split: ', stump.feature_index)
      print('Feature threshold: ', stump.threshold)
      print('Min error: ', stump.min_err)
```

	precision	recall	f1-score	support
-1	0.80	0.80	0.80	70
1	0.82	0.82	0.82	80
accuracy			0.81	150
macro avg	0.81	0.81	0.81	150
weighted avg	0.81	0.81	0.81	150

```
Feature to split: 13
Feature threshold: -0.24586126765035565
Min error: 0.14
```

```
[ ]:
```